

Introduction to Data Science

DSA1101

Semester 1, 2018/2019

Week 9

The Naïve Bayes Classifier

Naïve Bayes Classifier

- *Naïve Bayes* is a probabilistic classification method based on Bayes' theorem (or Bayes' law) with a few tweaks
- Bayes' theorem gives the relationship between the probabilities of two events and their conditional probabilities.

Naïve Bayes Classifier: Example

- The example is on using naïve Bayes classifier to predict whether employees would enroll in an onsite educational program based on feature variables such as Age, Income, JobSatisfaction and Desire.
- We will illustrate with both manual calculation and using the `naiveBayes` function in the [R](#) package 'e1071'.

Naïve Bayes Classifier: Example in R

- The CSV file 'sample1.csv' has been posted to IVLE, and contains 15 records.
- The last record does not contain any outcome label for Enrolls

```
1 > sample <- read.table("sample1.csv",header=TRUE,
2   sep=",")
3 > head(sample)
4   Age Income JobSatisfaction
5 1   <=30   High           No
6 2   <=30   High           No
7 3 31 to 40   High           No
8 4   >40 Medium           No
9 5   >40   Low           Yes
10 6   >40   Low           Yes
11   Desire Enrolls
12 1   Fair       No
13 2 Excellent    No
14 3   Fair      Yes
15 4   Fair      Yes
16 5   Fair      Yes
17 6 Excellent    No
```

Naïve Bayes Classifier: Example in R

- Two data frame objects called `traindata` and `testdata` are created for the naïve Bayes Classifier
- We will train the classifier using `traindata`, then make predictions for the single record in `testdata`

```
1 > traindata <- as.data.frame(sample[1:14,])
2 > testdata <- as.data.frame(sample[15,])
3 > testdata
4      Age Income JobSatisfaction Desire
5 15 <=30 Medium                Yes  Fair
6      Enrolls
7 15
```

Naïve Bayes Classifier: Example in R

- We will first illustrate the naïve Bayes classifier via manual computation
- Recall that we need to compute the probabilities $P(Y = y_j)$ for $j = 1, 2, \dots, k$.

```
1 > tprior <- table(traindata$Enrolls)
2 > tprior
3
4      No Yes
5  0    5  9
6 > tprior <- tprior/sum(tprior)
7 > tprior
8
9              No      Yes
10 0.0000000 0.3571429 0.6428571
```

Naïve Bayes Classifier: Example in R

- We will first illustrate the naïve Bayes classifier via manual computation
- Recall that we need to compute the probabilities $P(Y = y_j)$ for $j = 1, 2, \dots, k$.
- Here since Y is binary, we just compute $P(Y = 1)$ (for yes) and $P(Y = 0)$ (for no).

```
1 > tprior <- table(traindata$Enrolls)
2 > tprior
3
4      No Yes
5  0    5   9
6 > tprior <- tprior/sum(tprior)
7 > tprior
8
9              No      Yes
10 0.0000000 0.3571429 0.6428571
```


Naïve Bayes Classifier: Example in R

- Next, we need to compute the conditional probabilities $P(X_i = x_i | Y = 1)$ and $P(X_i = x_i | Y = 0)$, where $i = 1, 2, 3, 4$ for the feature variables $X = \{\text{Age}, \text{Income}, \text{JobSatisfaction}, \text{Desire}\}$.
- First, compute the conditional probabilities for Age:

```
1 > ageCounts <- table(traindata[,c("Enrolls", "Age")])
2 > ageCounts
3      Age
4 Enrolls  <=30  >40  31 to 40
5           0    0          0
6      No      3    2          0
7      Yes     2    3          4
8 > ageCounts <- ageCounts/rowSums(ageCounts)
9 > ageCounts
10      Age
11 Enrolls  <=30  >40  31 to 40
12
13      No  0.6000000  0.4000000  0.0000000
14      Yes 0.2222222  0.3333333  0.4444444
```

Naïve Bayes Classifier: Example in R

- We perform similar operations for Income, JobSatisfaction and Desire:

```
1 > incomeCounts <- table(traindata[,c("Enrolls", "
   Income")])
2 > incomeCounts <- incomeCounts/rowSums(
   incomeCounts)
3 > incomeCounts
4      Income
5 Enrolls      High      Low      Medium
6
7      No  0.4000000  0.2000000  0.4000000
8      Yes  0.2222222  0.3333333  0.4444444
```

Naïve Bayes Classifier: Example in R

- We perform similar operations for Income, JobSatisfaction and Desire:

```
1 > jsCounts <- table(traindata[,c("Enrolls", "
   JobSatisfaction")])
2 > jsCounts <- jsCounts/rowSums(jsCounts)
3 > jsCounts
4           JobSatisfaction
5 Enrolls           No           Yes
6
7      No  0.8000000  0.2000000
8      Yes 0.3333333  0.6666667
```

Naïve Bayes Classifier: Example in R

- We perform similar operations for Income, JobSatisfaction and Desire:

```
1 > desireCounts <- table(traindata[,c("Enrolls", "
  Desire")])
2 > desireCounts <- desireCounts/rowSums(
  desireCounts)
3 > desireCounts
4       Desire
5 Enrolls Excellent      Fair
6
7    No  0.6000000 0.4000000
8    Yes 0.3333333 0.6666667
```

Naïve Bayes Classifier: Example in R

- We finally compute the probability scores

$$P(Y = 1|X) \propto P(Y = 1) \times \prod_{i=1}^4 P(X_i = x_i|Y = 1)$$

and

$$P(Y = 0|X) \propto P(Y = 0) \times \prod_{i=1}^4 P(X_i = x_i|Y = 0)$$

Naïve Bayes Classifier: Example in R

- Computation of the probability scores:

```
1 prob_yes <-  
2 ageCounts["Yes",testdata[,c("Age")]]*  
3 incomeCounts["Yes",testdata[,c("Income")]]*  
4 jsCounts["Yes",testdata[,c("JobSatisfaction")]]*  
5 desireCounts["Yes",testdata[,c("Desire")]]*  
6 tprior["Yes"]  
7  
8  
9 prob_no <-  
10 ageCounts["No",testdata[,c("Age")]]*  
11 incomeCounts["No",testdata[,c("Income")]]*  
12 jsCounts["No",testdata[,c("JobSatisfaction")]]*  
13 desireCounts["No",testdata[,c("Desire")]]*  
14 tprior["No"]
```

Naïve Bayes Classifier: Example in R

- Since $P(Y = 1|X) > P(Y = 0|X)$, the predicted result for the test record is yes for Enrolls.

```
1 > prob_yes
2           Yes
3 0.02821869
4 > prob_no
5           No
6 0.006857143
```

Naïve Bayes Classifier: Example in R

- Alternatively, we can use the `naiveBayes` function in the R package 'e1071' to perform naïve Bayes classification:

```
1 > library(e1071)
2 >
3 > model <- naiveBayes(Enrolls ~
4 + Age+Income+JobSatisfaction+Desire,
5 + traindata, laplace=0)
6 >
7 > results <- predict(model, testdata)
8 > results
9 [1] Yes
10 Levels: No Yes
```


Naïve Bayes Classifier: Diagnostics

- Recall that for diagnostics of classifiers, we have learnt about the *confusion matrix* as well as measures such as accuracy, true positive rate etc.
- We will familiarize ourselves with one additional diagnostics tool, the Receiver Operating Characteristic (ROC) curve.

Naïve Bayes Classifier: Diagnostics

- Recall that the False Positive Rate (FPR) and True Positive Rate (TPR) are calculated as

$$\text{FPR} = \frac{FP}{FP + TN}$$

and

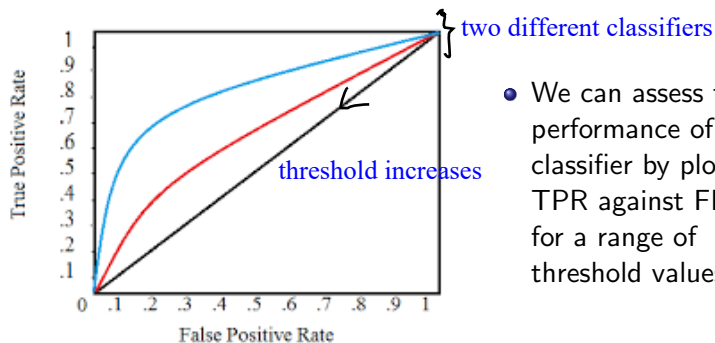
$$\text{TPR} = \frac{TP}{TP + FN}$$

		Predicted Class	
		Positive	Negative
	Actual Class		
	Positive	True Positives (TP)	False Negatives (FN)
	Negative	False Positives (FP)	True Negatives (TN)

Naïve Bayes Classifier: Diagnostics

- Recall that for classification using the majority rule, Y is predicted to be 1 if $\hat{Y} > 0.5$ and 0 otherwise.
- If the threshold is increased, then *less* test objects will be predicted to be 1, and so TP will be either constant or decreases. However, the sum $TP + FN$ is still constant because the number of objects with actual label $Y = 1$ is a constant in the test dataset, so TPR will either be constant or increases.
- Similarly, if the threshold is increased, FP will be either constant or increases, while the sum $FP + TN$ is a constant, so FPR will either be constant or increases.
- Therefore, in summary, if the threshold is increased, both TPR and FPR generally ~~increase~~ decrease
Find analogy in life?

Naïve Bayes Classifier: Diagnostics




- We can assess the performance of a classifier by plotting TPR against FPR for a range of threshold values

Naïve Bayes Classifier: Diagnostics

the blue curve is better,

higher area under the ROC curve,

two different classifiers



threshold increases

- A useful metric is to compute the area under the ROC curve (AUC)
- We will talk more about the ROC curve next week

for a given FPR, if $\Delta(\text{TPR})$ is maximum

-> the best threshold.

A good model should have a high TPR and low FPR.