

Association Rules

- 1 Introduction
- 2 *Apriori* Algorithm
- 3 Evaluation of Candidate Rules
- 4 Example in R
 - Groceries Data Set
 - Finding Frequent Itemsets
 - Generating Rules

1 Introduction

2 *Apriori* Algorithm

3 Evaluation of Candidate Rules

4 Example in R

- Groceries Data Set
- Finding Frequent Itemsets
- Generating Rules

Association Rules

- Association rules is a descriptive, not predictive, method often used to discover interesting relationships hidden in a large data set.
- The disclosed relationships can be represented as rules or frequent itemsets.
- Association rules are commonly used for mining transactions in databases.

Association Rules

- For example, given a large collection of retail transactions, in which each transaction consists of one or more items, association rules go through the items being purchased to see **what items are frequently bought together** and to discover a list of **rules that describe the purchasing behavior**.
- The goal with association rules is to discover interesting **relationships among the items**.
- The relationships that are interesting depend both on the business context and the nature of the algorithm being used for the discovery.

Association Rules



Source: *Data Science & Big Data Analytics*

1 Introduction

2 *Apriori* Algorithm

3 Evaluation of Candidate Rules

4 Example in R

- Groceries Data Set
- Finding Frequent Itemsets
- Generating Rules

New Term: “Itemset”

- In the example of a retail store, association rules are used over transactions that consist of one or more items.
- In fact, because of their popularity in mining customer transactions, association rules are sometimes referred to as **market basket analysis**.
- Each transaction can be viewed as the shopping basket of a customer that contains one or more items.
- This is also known as an **itemset**.
- The term *itemset* refers to a collection of items or individual entities that contain some kind of relationship.

Itemset

- An itemset containing k items is called a k -itemset.
- We will use the notation $\{\text{item 1, item 2, ..., item } k\}$ to denote a k -itemset.
- Computation of the association rules is typically based on itemsets.
- We will focus on the *Apriori* algorithm for generating association rules.

The Support of an Itemset

- One major component of *Apriori* algorithm is 'support'.
- Given an itemset L , the *support* of L is the percentage of transactions that contain L .
- For example, if 80% of all transactions contain itemset $\{bread\}$, then the support of $\{bread\}$ is 0.8.
- Similarly, if 60% of all transactions contain itemset $\{bread, butter\}$, then the support of $\{bread, butter\}$ is 0.6.


The Minimum Support

- A frequent itemset has items that appear together *often enough*.
- The term “often enough” is formally defined with a **minimum support** criterion.
- If the **minimum support is set at 0.5**, any itemset can be considered a frequent itemset if **at least 50% of the transactions contain this itemset**.
- For the previous example, both $\{bread\}$ and $\{bread, butter\}$ are considered frequent itemsets at the minimum support 0.5.
- If the minimum support is 0.7, only $\{bread\}$ is considered a frequent itemset.

Frequent Itemset

- If an itemset is considered **frequent**, then any subset of the frequent itemset must also be frequent.
- This is referred to as the *Apriori property* (or downward closure property).
- For example, if 60% of the transactions contain $\{bread, jam\}$, then at least 60% of all the transactions will contain $\{bread\}$ or $\{jam\}$.
- The *Apriori property* provides the basis for the *Apriori* algorithm.

Example



Source: *Data Science & Big Data Analytics*

Apriori Algorithm

- The Apriori algorithm takes a bottom-up iterative approach to uncovering the frequent itemsets by: **(1)** first determining all the possible items (or 1-itemsets, for example $\{bread\}$, $\{eggs\}$, $\{milk\}$,...) and **(2)** then identifying which among them are frequent based on a minimum support threshold (or the minimum support criterion).
- For example, when the minimum support threshold is set at 0.5, the algorithm identifies and retains those itemsets that appear in at least 50% of all transactions and discards (or “prunes away”) the itemsets that have a support less than 0.5 or appear in fewer than 50% of the transactions.

Apriori Algorithm

- In the next iteration of the Apriori algorithm, the identified frequent 1-itemsets are paired into 2-itemsets,
- For example, $\{bread, eggs\}$, $\{bread, milk\}$, $\{eggs, milk\}$, ... and again evaluated to identify which are the frequent 2-itemsets among them.
- This iterative process is repeated in the Apriori algorithm.
- At each iteration, the algorithm checks whether the support criterion can be met; if it can, the algorithm grows the itemset, repeating the process until it runs out of support or until the itemsets reach a predefined length.

Apriori Algorithm

- The growing and pruning process is repeated until no itemsets meet the minimum support threshold.
- Optionally, a threshold N can be set up to specify the maximum number of items the itemset can reach or the maximum number of iterations of the algorithm.
- Once completed, output of the *Apriori* algorithm is the collection of all the frequent k -itemsets.

Apriori Algorithm

- Finally, a collection of candidate rules is formed based on the frequent k -itemsets uncovered in the iterative process described earlier.
- For example, a frequent itemset $\{milk, eggs\}$ may suggest candidate rules $\{milk\} \rightarrow \{eggs\}$ and $\{eggs\} \rightarrow \{milk\}$.
- We can evaluate the appropriateness of these candidate rules using measures such as **confidence**, **lift**, and **leverage**.

- 1 Introduction
- 2 *Apriori* Algorithm
- 3 Evaluation of Candidate Rules
- 4 Example in R
 - Groceries Data Set
 - Finding Frequent Itemsets
 - Generating Rules

Confidence of a Rule

- *Confidence* is defined as the measure of certainty or trustworthiness associated with each discovered rule.
- Mathematically, the confidence for candidate rule $X \rightarrow Y$ is the percent of transactions that contain both X and Y out of all the transactions that contain X :

$$\text{Confidence}(X \rightarrow Y) = \frac{\text{Support}(X \wedge Y)}{\text{Support}(X)}$$

- Since $\text{Support}(X \wedge Y) \leq \text{Support}(X)$ always, we have $\text{Confidence}(X \rightarrow Y) \leq 1$.

Confidence of a Rule

- For example, if $\{bread, eggs, milk\}$ has a support of 0.15 and $\{bread, eggs\}$ also has a support of 0.15, the confidence of rule $\{bread, eggs\} \rightarrow \{milk\}$ is 1, which means 100% of the time a customer buys bread and eggs, milk is bought as well.
- The rule is therefore correct for 100% of the transactions containing bread and eggs.

Confidence of a Rule

- A relationship may be thought of as interesting when the algorithm identifies the relationship with a measure of confidence greater than or equal to a predefined threshold.
- This predefined threshold is called the **minimum confidence**.
- A higher confidence indicates that the rule ($X \rightarrow Y$) is more interesting or more trustworthy, based on the sample data set.

Confidence of a Rule

- Even though confidence can identify the interesting rules from all the candidate rules, it comes with a problem.
- Given rules in the form of $X \rightarrow Y$, confidence considers only the antecedent (X) and the co-occurrence of X and Y ; it does not take the consequent of the rule (Y) into concern.
- Other measures such as lift and leverage are designed to address this issue.

Lift of a Rule

- *Lift* measures how many times more often X and Y occur together than expected if they are statistically independent of each other.
- *Lift* is a measure of how X and Y are really related rather than coincidentally happening together:

$$Lift(X \rightarrow Y) = \frac{Support(X \wedge Y)}{Support(X) \times Support(Y)}.$$

- Lift is 1 if X and Y are statistically independent of each other.
- In contrast, a lift of $X \rightarrow Y$ greater than 1 indicates that there is some usefulness to the rule.

Lift of a Rule

- For example, assuming from a total of 1,000 transactions, with $\{milk, eggs\}$ appearing in 300 of them, $\{milk\}$ appearing in 500, and $\{eggs\}$ appearing in 400, then

$$Lift(milk \rightarrow eggs) = 0.3 / (0.5 \times 0.4) = 1.5.$$

- If $\{bread\}$ appears in 400 transactions and $\{milk, bread\}$ appears in 400, then

$$Lift(milk \rightarrow bread) = 0.4 / (0.5 \times 0.4) = 2.$$

Leverage of a Rule

- *Leverage* is a similar notion, but instead of using a ratio, leverage uses the difference:

$$\text{Leverage}(X \rightarrow Y) = \text{Support}(X \wedge Y) - \text{Support}(X) \times \text{Support}(Y)$$

- *Leverage* measures the difference in the probability of X and Y appearing together in the data set compared to what would be expected if X and Y were statistically independent of each other.
- In theory, leverage is 0 when X and Y are statistically independent of each other.
- A larger leverage absolute value indicates a stronger relationship between X and Y .

Leverage of a Rule

- For example, assuming a total of 1,000 transactions, with $\{milk, eggs\}$ appearing in 300 of them, $\{milk\}$ appearing in 500, and $\{eggs\}$ appearing in 400, then

$$Leverage(milk \rightarrow eggs) = 0.3 - (0.5 \times 0.4) = 0.1.$$

- If $\{bread\}$ appears in 400 transactions and $\{milk, bread\}$ appears in 400, then

$$Leverage(milk \rightarrow bread) = 0.4 - (0.5 \times 0.4) = 0.2.$$

Applications of Association Rules

- The term *market basket analysis* refers to a specific implementation of association rules mining that many companies use for a variety of purposes, including these:
 - 1 Broad-scale approaches to better merchandising-what products should be included in or excluded from the inventory each month
 - 2 Cross-merchandising between products and high-margin or high-ticket items
 - 3 Physical or logical placement of product within related categories of products
 - 4 Promotional programs-multiple product purchase incentives managed through a loyalty card program, etc.

Applications of Association Rules

- Besides market basket analysis, association rules are commonly used for recommender systems

Today's Recommendations For You

Here's a daily sample of items recommended for you. Click here to see more recommendations.



Applications of Association Rules

- Many online service providers such as Amazon and Netflix use recommender systems.
- Recommender systems can use association rules to discover related products or identify customers who have similar interests.
- For example, association rules may suggest that those customers who have bought product A have also bought product B .
- These findings provide opportunities for retailers to cross-sell their products.

Applications of Association Rules

- Clickstream analysis refers to the analytics on data related to web browsing and user clicks, which is stored on the client or the server side.
- Web usage log files generated on web servers contain huge amounts of information, and association rules can potentially give useful knowledge to web usage data analysts.
- For example, association rules may suggest that website visitors who land on page X click on links A , B , and C much more often than links D , E , and F .
- This observation provides valuable insight on how to better personalize and recommend the content to site visitors.

- 1 Introduction
- 2 *Apriori* Algorithm
- 3 Evaluation of Candidate Rules
- 4 Example in R
 - Groceries Data Set
 - Finding Frequent Itemsets
 - Generating Rules

- 1 Introduction
- 2 *Apriori* Algorithm
- 3 Evaluation of Candidate Rules
- 4 Example in R
 - Groceries Data Set
 - Finding Frequent Itemsets
 - Generating Rules

Grocery Store Transaction


- We will look at an example illustrating the application of the *Apriori algorithm* to grocery store transaction data.
- Using **R** and the 'arules' and 'arulesViz' packages, this example shows how to use the *Apriori algorithm* to generate frequent itemsets and rules and to evaluate and visualize the rules.

```
> #install.packages('arules')  
> #install.packages('arulesViz')  
> library('arules')  
> library('arulesViz')
```

Groceries Data

- The example uses the Groceries data set from the package 'arules' in R.¹
- The Groceries dataset is collected from 30 days of real-world point-of-sale transactions of a grocery store.
- The dataset contains 9835 transactions, and the items are aggregated into 169 categories.

```
> data(Groceries)
> Groceries
transactions in sparse format with
9835 transactions (rows) and
169 items (columns)
```

¹<https://cran.r-project.org/web/packages/arules/arules.pdf> 

Groceries Data

- The summary shows that the most frequent items in the dataset include items such as whole milk, other vegetables, rolls/buns, soda, and yogurt. These items are purchased more often than the others.

```
> summary(Groceries)
```

```
transactions as itemMatrix in sparse format with  
9835 rows (elements/itemsets/transactions) and  
169 columns (items) and a density of 0.02609146
```

most frequent items:

whole milk	other vegetables	rolls/buns	soda
2513	1903	1809	1776
yogurt	(Other)		
1372	34055		

element (itemset/transaction) length distribution:
sizes

1	2	3	4	5	6	7	8	9	10	11	12	13
2159	1643	1299	1005	855	645	545	438	350	246	182	117	78
17	18	19	20	21	22	23	24	26	27	28	29	32
29	14	14	9	11	4	6	1	1	1	1	3	1

Groceries Data

```
> inspect(head(Groceries))  
  items  
[1] {citrus fruit,  
    semi-finished bread,  
    margarine,  
    ready soups}  
[2] {tropical fruit,  
    yogurt,  
    coffee}  
[3] {whole milk}  
[4] {pip fruit,  
    yogurt,  
    cream cheese ,  
    meat spreads}  
[5] {other vegetables,  
    whole milk,  
    condensed milk,  
    long life bakery product}  
[6] {whole milk,  
    butter,
```

Groceries Data

- Data set 'Groceries' is a transaction class, as defined by the 'arules' package. A transactions class has three component slots:
 - 1 itemsetInfo: A data frame with vectors of the same length as the number of transactions
 - 2 itemInfo: A data frame to store item labels
 - 3 data: A binary incidence matrix that indicates which item labels appear in every transaction

Groceries Data

- `Groceries@itemInfo` display all 169 grocery labels as well as their categories.

```
> Groceries@itemInfo[1:10,]
```

	labels	level2	level1
1	frankfurter	sausage meat	and sausage
2	sausage	sausage meat	and sausage
3	liver loaf	sausage meat	and sausage
4	ham	sausage meat	and sausage
5	meat	sausage meat	and sausage
6	finished products	sausage meat	and sausage
7	organic	sausage	sausage meat and sausage
8	chicken	poultry meat	and sausage
9	turkey	poultry meat	and sausage
10	pork	pork meat	and sausage

Groceries Data

- `Groceries@data` indicates which item labels appear in every transaction.
- `|` indicates that the item appears in transaction, and `.` otherwise.

```
> Groceries@data[,100:110]
```

```
169 x 11 sparse Matrix of class "ngCMatrix"
```

```
[1,] . . . | . . . . . . .  
[2,] . . | . | . . . . .  
[3,] . . . . . . . . . .  
[4,] . . . . . . . . . .  
[5,] . . . . . . . . . .  
[6,] . . . . . . . . . .  
[7,] . . . . . . . . . .  
[8,] . . . . . . . . . .  
[9,] . . . . . . . . . .  
[10,] . . | . . . . . . .  
[11,] . . . . . . . | . . .  
[12,] . . . . . . . . . .  
[13,] . . . . . . . . . .  
[14,] | . . . . . . | . . .  
[15,] | . . . . . . . . . |
```

Groceries Data

- The following code displays the 1st to 5th transactions of the Groceries data set, similar as from slide 36.
- `[,1 : 5]` can be changed to `[,1 : 9835]` to display all the transactions.

```
> apply(Groceries@data[,1:5], 2,  
+       function(r) paste(Groceries@itemInfo[r,"labels"],  
+       collapse=", "))  
[1] "citrus fruit, semi-finished bread, margarine, ready soups"  
[2] "tropical fruit, yogurt, coffee"  
[3] "whole milk"  
[4] "pip fruit, yogurt, cream cheese , meat spreads"  
[5] "other vegetables, whole milk, condensed milk, long life bak
```


- 1 Introduction
- 2 *Apriori* Algorithm
- 3 Evaluation of Candidate Rules
- 4 Example in R
 - Groceries Data Set
 - Finding Frequent Itemsets
 - Generating Rules

apriori() Function

- The `apriori()` function from the 'arule' package implements the *Apriori algorithm* to create frequent itemsets.
- Note that, by default, the `apriori()` function executes all the iterations at once.
- However, to illustrate how the *Apriori algorithm* works, the code examples in this section manually set the parameters of the `apriori()` function to simulate each iteration of the algorithm.

apriori() Function

- Assume that the **minimum support threshold** is set to **0.02** based on management discretion.
- Because the data set contains 9,835 transactions, an itemset should appear 196-197 times to be considered a frequent itemset.
- The first iteration of the *Apriori algorithm* computes the support of each product in the data set and retains those products that satisfy the minimum support.

apriori() Function

- The following code identifies 59 frequent 1-itemsets that satisfy the minimum support.
- The parameters of **apriori()** specify the minimum and maximum lengths of the itemsets, the minimum support threshold, and the target indicating the type of association mined.

apriori() to get frequent 1-itemsets

```
> itemsets.1 <- apriori(Groceries, parameter=list(minlen=1,  
+          maxlen=1,support=0.02, target="frequent itemsets"))
```

Apriori

Parameter specification:

confidence	minval	smax	arem	aval	originalSupport	maxtime	support
NA	0.1	1	none	FALSE	TRUE	5	0.02
maxlen	target	ext					
1	frequent itemsets	TRUE					

Algorithmic control:

filter	tree	heap	memopt	load	sort	verbose
0.1	TRUE	TRUE	FALSE	TRUE	2	TRUE

Absolute minimum support count: 196

set item appearances ...[0 item(s)] done [0.00s].

set transactions ...[169 item(s), 9835 transaction(s)] done [0.00s].

sorting and recoding items ... [59 item(s)] done [0.00s].

creating transaction tree ... done [0.00s].

```
> summary(itemsets.1)
```

```
set of 59 itemsets
```

```
most frequent items:
```

frankfurter	sausage	ham	meat	chicken	(Oth
1	1	1	1	1	

```
element (itemset/transaction) length distribution:sizes
```

```
1
```

```
59
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
1	1	1	1	1	1

```
summary of quality measures:
```

support	count
Min. :0.02105	Min. : 207.0
1st Qu.:0.03015	1st Qu.: 296.5
Median :0.04809	Median : 473.0
Mean :0.06200	Mean : 609.8
3rd Qu.:0.07666	3rd Qu.: 754.0
Max. :0.25552	Max. :2513.0

Most 10 Frequent 1-Itemsets

The following code uses the `inspect()` function to display the top 10 frequent 1-itemsets sorted by **their support**.

```
> inspect(head(sort(itemsets.1, by = "support"), 10))
```

	items	support	count
[1]	{whole milk}	0.25551601	2513
[2]	{other vegetables}	0.19349263	1903
[3]	{rolls/buns}	0.18393493	1809
[4]	{soda}	0.17437722	1715
[5]	{yogurt}	0.13950178	1372
[6]	{bottled water}	0.11052364	1087
[7]	{root vegetables}	0.10899847	1072
[8]	{tropical fruit}	0.10493137	1032
[9]	{shopping bags}	0.09852567	969
[10]	{sausage}	0.09395018	924

apriori() to get frequent 2-itemsets

- In the next iteration, the list of frequent 1-itemsets is joined onto itself to form all possible candidate 2-itemsets.
- For example, 1-itemsets $\{whole\ milk\}$ and $\{soda\}$ would be joined to become a 2-itemset $\{wholemilk, soda\}$.
- The algorithm computes the support of each candidate 2-itemset and retains those that satisfy the minimum support.
- The output that follows shows that 61 frequent 2-itemsets have been identified.


```
> itemsets.2 <- apriori(Groceries, parameter=list(minlen=2,
+           maxlen=2, support=0.02, target="frequent itemsets"))
Apriori
```

Parameter specification:

confidence	minval	smax	arem	aval	originalSupport	maxtime	support
NA	0.1	1	none	FALSE	TRUE	5	0.02
maxlen	target	ext					
2	frequent itemsets	TRUE					

Algorithmic control:

filter	tree	heap	memopt	load	sort	verbose
0.1	TRUE	TRUE	FALSE	TRUE	2	TRUE

Absolute minimum support count: 196

```
set item appearances ...[0 item(s)] done [0.00s].
set transactions ...[169 item(s), 9835 transaction(s)] done [0.00s].
sorting and recoding items ... [59 item(s)] done [0.00s].
creating transaction tree ... done [0.00s].
checking subsets of size 1 2 done [0.00s].
sorting transactions ... done [0.00s].
```

```
> summary(itemsets.2)
```

```
set of 61 itemsets
```

```
most frequent items:
```

whole milk	other vegetables	yogurt	rolls/buns
25	17	9	9
soda	(Other)		
9	53		

```
element (itemset/transaction) length distribution:sizes
```

```
2
```

```
61
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
2	2	2	2	2	2

```
summary of quality measures:
```

support	count
Min. :0.02003	Min. :197.0
1st Qu.:0.02227	1st Qu.:219.0
Median :0.02613	Median :257.0
Mean :0.02951	Mean :290.3

Most 10 Frequent 2-Itemsets

The top 10 frequent 2-itemsets sorted by **their support**.

```
> inspect(head(sort(itemsets.1, by = "support"), 10))
```

	items	support	count
[1]	{whole milk}	0.25551601	2513
[2]	{other vegetables}	0.19349263	1903
[3]	{rolls/buns}	0.18393493	1809
[4]	{soda}	0.17437722	1715
[5]	{yogurt}	0.13950178	1372
[6]	{bottled water}	0.11052364	1087
[7]	{root vegetables}	0.10899847	1072
[8]	{tropical fruit}	0.10493137	1032
[9]	{shopping bags}	0.09852567	969
[10]	{sausage}	0.09395018	924

- Notice that whole milk appears six times in the top 10 2-itemsets ranked by support.
- As seen earlier, whole milk has the highest support among all the 1-itemsets.
- These top 10 2-itemsets with the highest support may not be interesting; this highlights the limitations of using support alone.

apriori() to get frequent 3-itemsets

- Next, the list of frequent 2-itemsets is joined onto itself to form candidate 3-itemsets.
- For example $\{other\ vegetables, wholemilk\}$ and $\{whole\ milk, rolls/\ buns\}$ would be joined as $\{other\ vegetables, whole\ milk, rolls/\ buns\}$.
- The algorithm retains those itemsets that satisfy the minimum support.
- The following output shows that only **two** frequent 3-itemsets have been identified.

Frequent 3-Itemsets

```
> itemsets.3 <- apriori(Groceries, parameter=list(minlen=3,  
+      maxlen=3, support=0.02, target="frequent itemsets"))  
Apriori
```

Parameter specification:

confidence	minval	smax	arem	aval	originalSupport	maxtime	support
NA	0.1	1	none	FALSE	TRUE	5	0.02
maxlen	target	ext					
3	frequent itemsets	TRUE					

Algorithmic control:

filter	tree	heap	memopt	load	sort	verbose
0.1	TRUE	TRUE	FALSE	TRUE	2	TRUE

Absolute minimum support count: 196

set item appearances ...[0 item(s)] done [0.00s].

set transactions ...[169 item(s), 9835 transaction(s)] done [0.00s].

sorting and recoding items ... [59 item(s)] done [0.00s].

creating transaction tree ... done [0.00s].

Frequent 3-Itemsets

There are only **two** frequent 3-itemsets:

```
> inspect(sort(itemsets.3, by = "support"))
```

	items	support	count
[1]	{root vegetables, other vegetables, whole milk}	0.02318251	228
[2]	{other vegetables, whole milk, yogurt}	0.02226741	219

Frequent 4-Itemsets

```
> itemsets.4 <- apriori(Groceries, parameter=list(minlen=4,  
+      maxlen=4,support=0.02, target="frequent itemsets"))
```

Apriori

Parameter specification:

confidence	minval	smax	arem	aval	originalSupport	maxtime	support	maxlen	target	ext
NA	0.1	1	none	FALSE	TRUE	5	0.02	4	frequent itemsets	TRUE

Algorithmic control:

filter	tree	heap	memopt	load	sort	verbose
0.1	TRUE	TRUE	FALSE	TRUE	2	TRUE

Absolute minimum support count: 196

set item appearances ...[0 item(s)] done [0.00s].

set transactions ...[169 item(s), 9835 transaction(s)] done [0.00s].

sorting and recoding items ... [59 item(s)] done [0.00s].

creating transaction tree ... done [0.00s].

Frequent 4-Itemsets

```
> summary(itemsets.4)
set of 0 itemsets
```

- With a minimum support of 0.02, there is NO frequent 4-itemset found.
- The iterations run out of support when $k = 4$.
- Therefore, the frequent itemsets contain: 59 frequent 1-itemsets, 61 frequent 2-itemsets, and 2 frequent 3-itemsets.
- When the `maxlen` parameter is not set, the algorithm continues each iteration until it runs out of support or until k reaches the default `maxlen=10`.

- 1 Introduction
- 2 *Apriori* Algorithm
- 3 Evaluation of Candidate Rules
- 4 Example in R
 - Groceries Data Set
 - Finding Frequent Itemsets
 - Generating Rules

Generating Rules

- The `apriori()` function can also be used to generate rules.
- Assume that the minimum support threshold is now set to a lower value 0.001, and the minimum confidence threshold is set to 0.6.
- A lower minimum support threshold allows more rules to show up.
- The following code creates 2,918 rules from all the transactions in the Groceries data set that satisfy both the minimum support and the minimum confidence.

```
> rules <- apriori(Groceries, parameter=list(support=0.001,  
+ confidence=0.6, target = "rules"))
```

Apriori

Parameter specification:

confidence	minval	smax	arem	aval	originalSupport	maxtime	support
0.6	0.1	1	none	FALSE	TRUE	5	0.001

maxlen	target	ext
10	rules	TRUE

Algorithmic control:

filter	tree	heap	memopt	load	sort	verbose
0.1	TRUE	TRUE	FALSE	TRUE	2	TRUE

Absolute minimum support count: 9


```
set item appearances ...[0 item(s)] done [0.00s].  
set transactions ...[169 item(s), 9835 transaction(s)] done [0.00s].  
sorting and recoding items ... [157 item(s)] done [0.00s].  
creating transaction tree ... done [0.00s].  
checking subsets of size 1 2 3 4 5 6 done [0.01s].  
writing ... [2918 rule(s)] done [0.00s].
```

Visualizing the Rules

- The command `plot(rules)` display the scatter plot of all the 2,918 rules, where the X-axis is the support, the Y-axis is the confidence, and the shading is the lift.
- The scatter plot shows that, of the 2,918 rules generated from the Groceries data set, the highest lift occurs at a low support and a low confidence.

```
> plot(rules, measure = c("support", "confidence"),  
+       shading = "lift", col = "black")#, limit = 100)
```

Scatter plot for 2918 rules



Some Top Rules

- The `inspect()` function can display content of the rules generated previously.
- The following code shows the top three rules sorted by the lift. Rule $\{Instant\ food\ products, soda\} \rightarrow \{hamburger\ meat\}$ has the highest lift of ≈ 19 .

```
> inspect(head(sort(rules, by="lift"), 3))
```

	lhs	rhs	support
[1]	{Instant food products, soda}	=> {hamburger meat}	0.00122013
[2]	{soda, popcorn}	=> {salty snack}	0.00122013
[3]	{ham, processed cheese}	=> {white bread}	0.00193187

	coverage	lift	count
[1]	0.001931876	18.99565	12
[2]	0.001931876	16.69779	12
[3]	0.003050330	15.04549	19

Top 5 Rules with Highest Lift

- We can also visualize the top five rules with the highest lift using the following code.
- ```
> highLiftRules <- head(sort(rules, by="lift"), 5)
> plot(highLiftRules, method = "graph", engine = "igraph",
+ edgeCol = "blue", alpha = 1)
```
- The size of the node is sorted by the support.
- The darkness of the red color represents the change in lift, darker means higher lift.



## Graph for 5 rules

size: support (0.001 - 0.002)  
color: lift (11.279 - 18.996)

